

Table of Contents

- [Executive Summary](#)
- [Architecture Overview Diagram](#)
- [End-to-End Data Flow \(Step-by-Step\)](#)
- [Database Retrieval Layer](#)
- [Context & Prompt Assembly](#)
- [Tool-Calling Mechanism](#)
- [Agent-Module Interaction Matrix](#)
- [Memory & Conversation System](#)
- [Output Guardrails & Token Governance](#)
- [MySQL Schema Reference \(AI Tables\)](#)

Executive Summary

StratRoom's chat engine is a **tool-calling agent architecture**, not a vector-based RAG system. There are no embeddings, no vector database (e.g. Chroma/FAISS/Pinecone), and no LangChain. Instead:

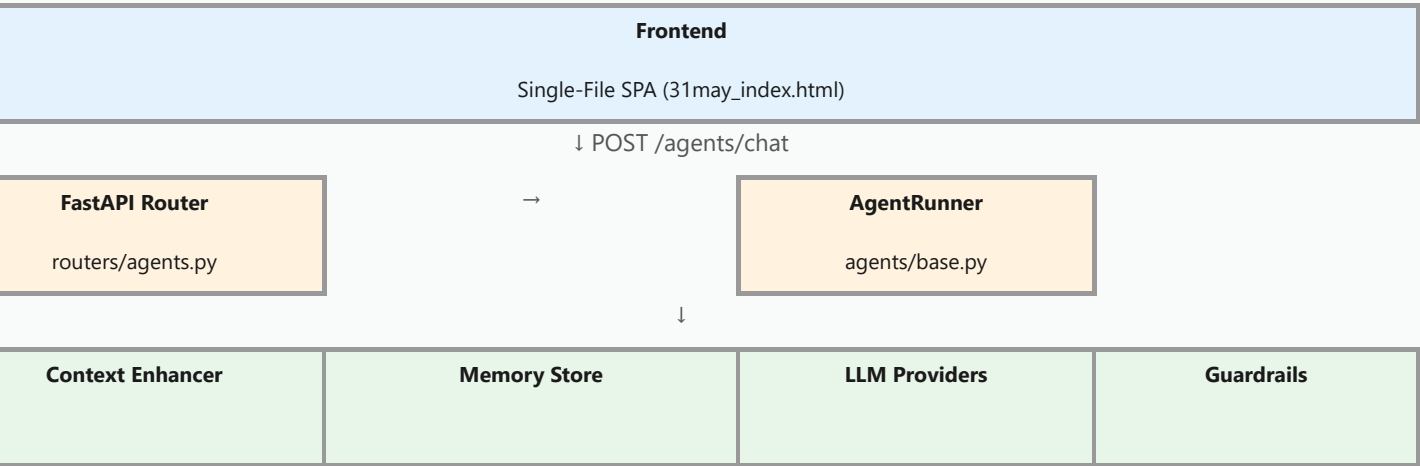
- The **LLM** receives a system prompt containing module-specific instructions and live organizational data.
- The LLM issues `[TOOL_CALL:tool_name:args]` markers in its response.
- The **server** parses these markers, executes MySQL queries via tool functions, and feeds the results back to the LLM.
- The LLM then produces a natural-language summary of the tool results for the user.

Long-term context is provided by a MySQL-backed memory system (`ai_memory` table) that stores confidence-scored insights from past conversations, serving a similar role to vector retrieval but via SQL rather than similarity search.

Architecture Overview Diagram

High-Level System Architecture

Figure 1 — High-Level System Architecture




```
ai_memory
org_id = %s AND agent_name = %s
user_id IS NULL OR user_id != %s)
BY confidence DESC, accessed_at DESC LIMIT %s
```

up to 8 memories ranked by confidence then recency.

B — Recent conversation summaries ([ai/query_enhancer.py:64-103](#)):

```
Get up to 3 most recent conversations for this agent+user
SELECT c.id, c.title, c.created_at
FROM agent_conversations c
WHERE c.agent_name = %s AND c.user_id = %s
ORDER BY c.created_at DESC LIMIT 3

For each conversation, get last 3 messages
SELECT role, content FROM agent_messages
WHERE conversation_id = %s ORDER BY id DESC LIMIT 3
```

Messages are truncated to 150 chars each and formatted as **User: ... | Agent:**

C — Org context hint:

```
ORGANIZATION ---
Organization ID: {org_id} | Agent: {agent_name}
```

D — System Prompt Assembly

[agents/base.py:73-76](#)

The system prompt is assembled in three parts:

```
PROMPTS[agent_name]}          ← Step 4a: Agent-specific persona + tool docs
enhanced}                    ← Step 4b: Memories + recent convos + org hint
--- CURRENT ORGANIZATION DATA ---\n{context} ← Step 4c: Live module data
```

Loads from **AGENT_PROMPTS** ([agents/prompts.py:1-346](#)), which defines 12 agent personas. Each prompt includes: - Role description - Available tools with **[TOOL_CALL:...]** syntax documentation - Response formatting instructions

Truncates context to **MAX_PROMPT_LENGTH** (default: **50,000** characters) if it exceeds the limit ([base.py:57-58](#)).


E — Conversation History Load


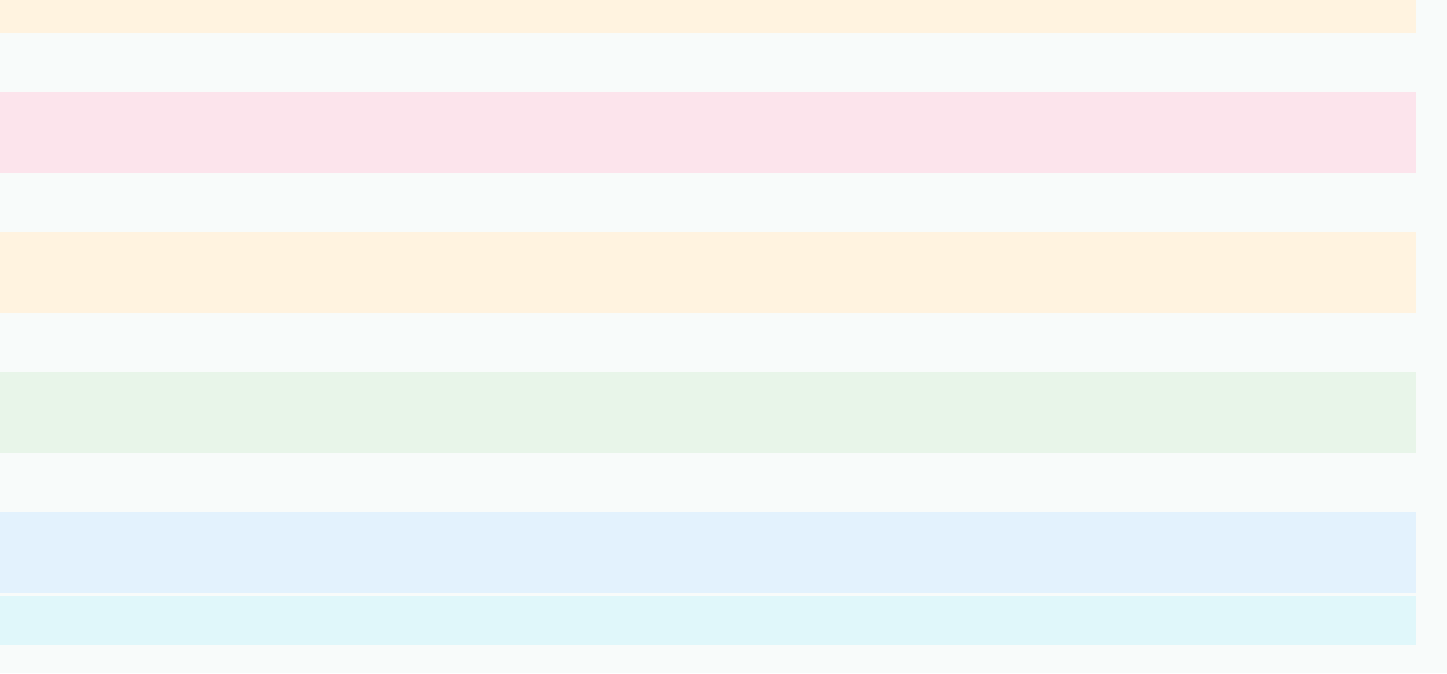
[agents/base.py:79-81](#)



```
conversation_id:
history = await self._load_history(conversation_id)
messages = history
```

Fetches the last **20 messages** from **agent_messages** via the Java DB service HTTP endpoint ([base.py:729-735](#)). This gives the LLM continuity across turns.

F — Intent Auto-Detection & Query Rewriting







```

ManagementList → SELECT ... FROM audit_management
ieveTaskList/ → SELECT ... FROM task_details
ingManagementList/ → SELECT ... FROM meeting_management
liance → SELECT ... FROM compliance_details
eCardList → SELECT ... FROM score_card
ByUser → SELECT ... FROM employee_details WHERE email = ?
oyeeDetailsList → SELECT ... FROM employee_details
List → SELECT ... FROM swot_details
ellList → SELECT ... FROM pestel_details

```

own paths fall through to HTTP calls to the Java microservices (ports 9010-9060).

SAC Filtering Patterns

functions implement **role-based access control**:

users see all data across orgs (unscoped queries):

```

min task query (no WHERE clause on org)
T t.ID, t.task_value, t.owner, t.priority, t.status
task_details t
BY FIELD(t.priority, 'Critical', 'High', 'Medium', 'Low'), t.ID

```

for users are scoped via `employee_details` JOIN:

```

mber task query (org + ownership scoped)
T t.ID, t.task_value, t.owner, t.priority, t.status
task_details t
employee_details e ON e.emp_id = t.owner
e.org_id = %s AND t.owner = %s
BY FIELD(t.priority, 'Critical', 'High', 'Medium', 'Low'), t.ID

```

/ resolution always starts with email → emp_id lookup:

```

T emp_id, org_id FROM employee_details
LOWER(email_address) = LOWER(%s) LIMIT 1

```

rd RBAC adds an additional `assigned_user_id` filter for non-admin, non-manager users (`scorecard_tools.py:60-61`).

ON Blob Pattern

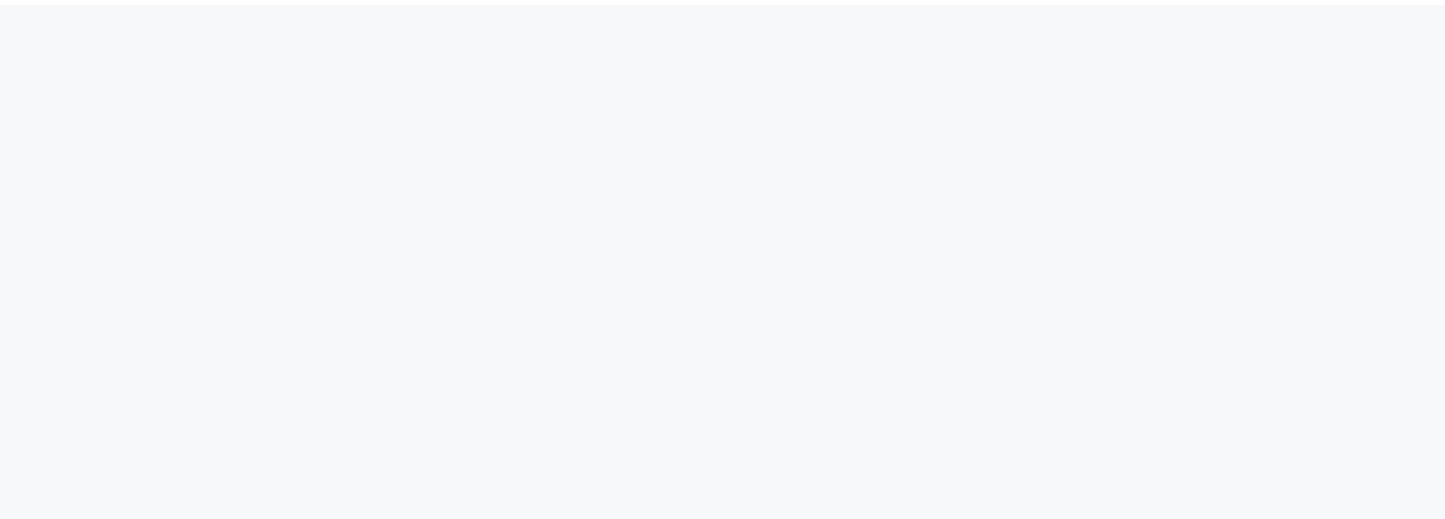
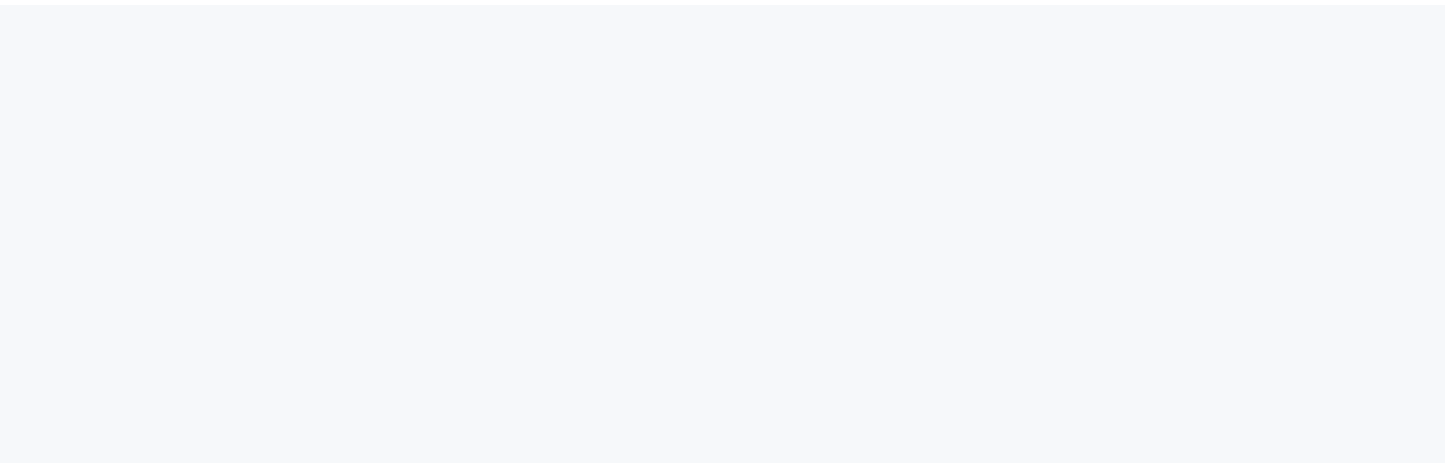
MySQL tables store structured data in a JSON column (e.g., `task_value`, `risk_value`, `initiative_value`, `incident_value`, `n_value`). The bridge parses these via `bridge._parse_json_col()`. When updating, the full JSON blob is read, merged with changes, then back.

Context & Prompt Assembly

Agent System Prompts

`ents/prompts.py:1-346`

personas are defined in the `AGENT_PROMPTS` dictionary:



```
CURRENT ORGANIZATION DATA ===
RISK AGENT DATA ===

RISKS ---
| risk_value={"name":"Cyber Breach","score":"20",...} | owner=101 | status=APPROVED
| risk_value={"name":"Vendor Lock-in","score":"15",...} | owner=105 | status=APPROVED

AUDIT ---
| managementvalue={...} | owner=101 | active=1

COMPLIANCE ---
| complain_value={...} | status=Compliant | risklevel=Low
```

Token & Character Limits

re/config.py:46-56

Parameter	Default	Purpose
CHAT_INPUT_CHARS	8,000	Max user message length
PROMPT_LENGTH	50,000	Max total system prompt (context block)
RESPONSE_CHARS	16,000	Hard output cap on LLM response
LLM_MAX_TOKENS	2,048	Provider <code>max_tokens</code> ceiling
TOOL_RESULT_CHARS	4,000	Max tool result fed back to LLM

Tool-Calling Mechanism

Tool Call Protocol

LLM emits tool calls as inline text markers:

```
_CALL:query_risks:min_heat=10]
_CALL:create_task:title=Deploy MFA,owner=dominic@demo.com,priority=High]
_CALL:update_task_status:id=102,status=completed]
```

```
: [TOOL_CALL:{tool_name}:{key1}={value1},{key2}={value2}]
```

Regex parser (base.py:31):

```
_CALL_RE = re.compile(r"\[TOOL_CALL:([^\]:]+)(?:([^\]]*))?\]")
```

Group 1: tool name (e.g., `query_risks`)

Group 2: optional args string (e.g., `min_heat=10`)

Argument Parsing

base.py:395-412 (`_parse_kv`)

```
        pass
    result[k.strip()] = v.strip() if isinstance(v, str) else v
return result
```

Complete Tool Inventory

Name	File	MySQL Table	Operation
task_details	task_tools.py:87	task_details	SELECT (with RBAC)
update_task_status	task_tools.py:276	task_details	UPDATE (status column + JSON)
update_task_progress	task_tools.py:424	task_details	UPDATE (progress in JSON, auto-status)
insert_task	task_tools.py:574	task_details	INSERT
insert_risk_mitigation_task	task_tools.py:698	task_details + risk_details	INSERT (linked task)
get_risks	risk_tools.py:22	risk_details	SELECT (with RBAC)
insert_risk	risk_tools.py:146	risk_details	INSERT
update_risk	risk_tools.py:299	risk_details	UPDATE (JSON blob merge)
delete_risk	risk_tools.py:497	risk_details	DELETE (admin only)
run_simulation	risk_tools.py:588	risk_details	SELECT → Monte Carlo simulation
get_scorecards	scorecard_tools.py:46	scorecard_kpis	SELECT (deduped, junk-filtered)
get_scorecard_summary	scorecard_tools.py:111	scorecard_kpis	SELECT (aggregated)
insert_scorecard	scorecard_tools.py:164	scorecard_kpis	INSERT
update_scorecard	scorecard_tools.py:219	scorecard_kpis	UPDATE
delete_scorecard	scorecard_tools.py:304	scorecard_kpis	DELETE (admin only)
get_initiatives	initiative_tools.py:47	initiatives_details	SELECT (with RBAC)
update_initiative_progress	initiative_tools.py:114	initiatives_details	UPDATE (JSON blob)
update_incident_status	incident_tools.py:44	universal_incident	UPDATE (nested JSON)
get_decisions	decision_tools.py:60	decisions	SELECT (with RBAC)
insert_decision	decision_tools.py:125	decisions	INSERT
update_decision_status	decision_tools.py:219	decisions	UPDATE (column + JSON)


ML Prediction Tools (Registry)

ml_tool_registry.py:52-172

Post-based prediction models are registered but invoked via a separate ML tool registry (not through the `[TOOL_CALL:...]` mechanism):

	Model File	Description
predict_risk	risk_model.joblib	Residual risk score prediction
predict_revenue	revenue_model.joblib	Revenue attainment %
predict_attrition	attrition_model.joblib	Employee attrition probability

--	--	--	--



	If present: load last 20 messages from <code>agent_messages</code> via <code>_load_history()</code>
	If absent: start a fresh message list
	Append the new user message and send the assembled <code>messages</code> to the LLM
	After the LLM responds, create a conversation via <code>POST /conversations</code> if none existed
	Save both user and assistant messages to <code>agent_messages</code>
	Store an insight in <code>ai_memory</code> , then prune stale memories

MySQL Tables

	Purpose	Key Columns
<code>_conversations</code>	Chat session metadata	<code>id</code> , <code>agent_name</code> , <code>user_id</code> , <code>org_id</code> , <code>title</code> , <code>created_at</code>
<code>_messages</code>	Individual chat messages	<code>id</code> , <code>conversation_id</code> , <code>role</code> , <code>content</code> , <code>created_at</code>
<code>_memory</code>	Long-term insight storage	<code>id</code> , <code>user_id</code> , <code>org_id</code> , <code>agent_name</code> , <code>insight</code> , <code>confidence</code> , <code>access_count</code>
<code>_agent_runs</code>	Per-call observability	<code>id</code> , <code>agent_name</code> , <code>provider</code> , <code>model</code> , <code>status</code> , <code>duration_ms</code> , <code>token_count</code>

Output Guardrails & Token Governance

Output Sanitization Pipeline

Figure 4 — Memory Lifecycle

	Store: LLM response passed to <code>store_memory()</code> (ai/memory.py)
	Filtered: trivial exchanges (greetings/small talk), responses < 40 chars, messages < 10 chars → skipped
	Confidence scored 0.0-1.0 (base 0.4 + data-richness bonuses); rejected if < 0.3
	Deduplicated via first-80-chars comparison against existing <code>ai_memory</code> rows
	INSERT into <code>ai_memory</code> (<code>user_id</code> , <code>org_id</code> , <code>agent_name</code> , <code>insight</code> , <code>confidence</code> , <code>source</code>)
	Retrieve: SELECT user-specific memories ordered by confidence DESC, accessed_at DESC (limit 8)
	Fallback: if fewer than 3 user-specific, supplement with org-wide memories
	Prune: after each store, if count > 200 per agent, delete lowest access_count + confidence + oldest

Output Validation Pipeline

Figure 7 — Output Validation Pipeline

	Raw LLM output from provider
	Empty response (< 5 chars) → safe fallback message returned
	Length > 50,000 chars → truncated with marker

